Interfaculty
Center
Data processing
Statistics

Sample size calculation with GPower

Wilfried Cools (ICDS) & Sven Van Laere (BiSI)
https://www.icds.be/

sample size calculation: why & how

- what about my sample size ?
- more is better ≡ observation → information
  - but increasingly less so
  - but limited resources (time/money)
  - but ethical and practical considerations
- how many is good enough ?
  - depends on what you want
- workshop → how to calculate it ?
  - understand the reasoning
  - apply for simple but common situations
- not one simple formula for all → GPower to the rescue

sample size calculation: a design issue

- linked to statistical inference
  - `testing` → power [probability to detect existing `effects`]
  - `estimation` → accuracy [size of `confidence intervals`]
- before data collection, during design of study
  - requires understanding: future data, analysis, inference (effect size, focus, ...)
  - conditional on assumptions & decisions
- not always possible nor meaningful !
  - easier for experiments (control), less for observational studies
  - easier for confirmatory studies, much less for exploratory studies
  - NO retrospective power analyses → OK for future study only

Hoenig, J., & Heisey, D. (2001). The Abuse of Power: The Pervasive Fallacy of Power Calculations for Data Analysis. *The American Statistician, 55*, 19–24.

simple example

- evaluation of radiotherapy to reduce a tumor in mice
- comparing treatment group with control (=conditions)
- tumor induced, random assignment treatment or control (equal if no effect)
- after 20 days, measurement of tumor size (=observations)
- analysis:
  - unpaired t-test to compare averages for treatment and control
- goal:
  - if the average size in treatment is at least 20% less than control
    then we want to detect it (significance)
- the main issue:
  - how to calculate the required sample size to determine the effect aimed for ?

overview

- PART I: building blocks in action for t-test
  - sizes: effect size, sample size
  - errors: type I (α), type II (β)
  - distributions: Ho, Ha
  - criterion: confidence (estimation), power (testing)
- PART II: moving beyond independent t-test
  - dependent groups
  - non-parametric distributions
  - multiple groups (ANOVA: omnibus, pairwise, focused)
  - proportions, correlations, ...

PART I: building blocks in action for t-test

reference example

- sample sizes easy and meaningful to calculate for well understood problems
- apriori specifications
  - intend to perform a statistical test
  - comparing 2 equally sized groups
  - to detect difference of at least 2
  - assuming an uncertainty of 4 SD on each mean
  - which results in an effect size of .5
  - evaluated on a Student t-distribution
  - allowing for a type I error prob. of .05 (α)
  - allowing for a type II error prob. of .2 (β)
- sample size conditional on specifications being true

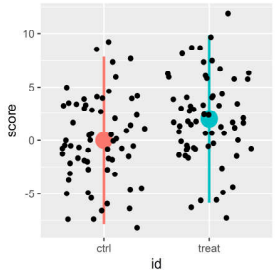

https://icds.shinyapps.io/shinyt/

a formula you could use

- for this particular case:
  - sample size (n → ?)
  - difference (d=signal → 2)
  - uncertainty (σ=noise → 4)
  - type I errors (α → .05, so $Z_{\alpha/2}$ → -1.96)
  - type II errors (β → .2, so Z_{β} → -0.84)

$$n = \frac{(Z_{\alpha/2}+Z_{\beta})^2 * 2 * \sigma^2}{d^2}$$

$$n = \frac{(-1.96-0.84)^2 * 2 * 4^2}{2^2} = 62.79$$

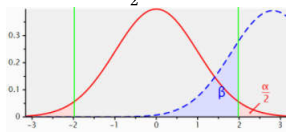

- sample size = 2 groups x 63 observations = 126
- note: formula's are test and statistic specific but logic remains same
- this and other formula's implemented in various tools, our focus: `GPower`

GPower: a useful tool

- popular and well established
- free @ http://www.gpower.hhu.de/
- implements wide variety of tests
- implements various visualizations
- documented fairly well
- note: not all tests are included !

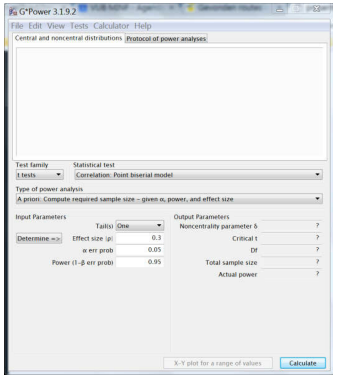

GPower: the building blocks in action

- For this example, calculate sample size based on
  - effect size (difference of interest, scaled on standard deviation)
  - type I and type II error

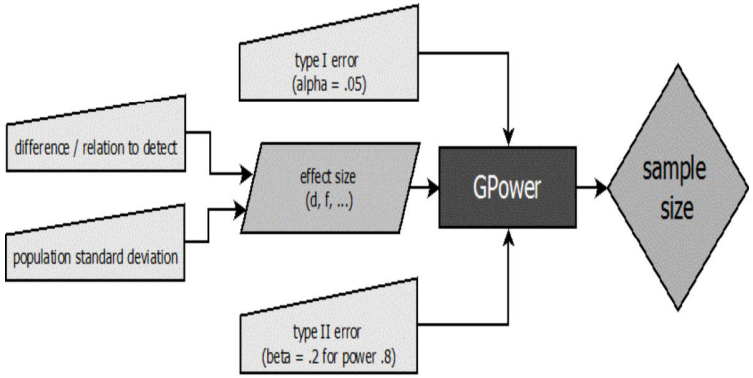

GPower input

- t-test : difference two indep. means
- apriori: calculate sample size
- effect size = standardized difference
  - Cohen's d
  - d = |difference| / SD_pooled
  - d = |`0-2`| / `4` = `.5`
- α = `.05`, 2 - tailed ($\alpha/2 \rightarrow$ .025 & .975)
- $power = 1 - \beta$ = `.8`
- allocation ratio = `1`
- `~ reference example`

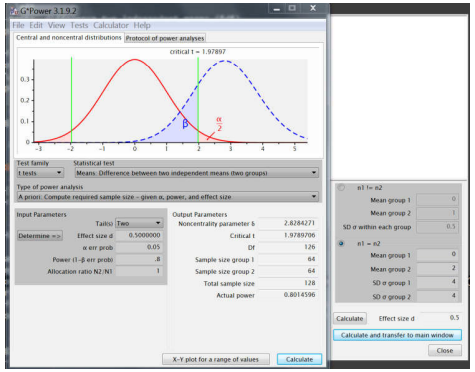

GPower output

- sample size (n) = 64 x 2 = (`128`)
- degrees of freedom (df) = 126 (128 - 2)
- plot showing null `Ho` and alternative `Ha` distribution
  - in GPower central and non-central distribution
  - `Ho` & critical value → decision boundaries
    - critical t = 1.979, `qt(.975,126)`
  - `Ha`, shift with non-centrality parameter → truth
    - non centrality parameter (δ) = 2.8284
      `2/(4*sqrt(2))*sqrt(64)`
- power ≥ .80 (1-β) = 0.8015

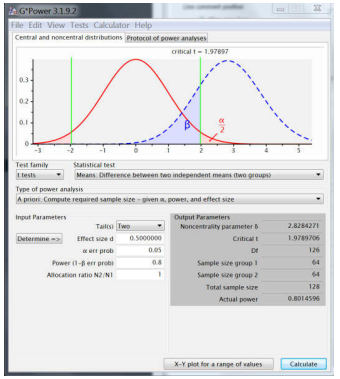

reference example protocol

t tests - Means: Difference between two independent means (two groups)
Analysis: A priori: Compute required sample size

Input: Tail(s) = Two
Effect size `d` = `0.5000000`
α err prob = `0.05`
Power (`1-β` err prob) = `.8`
Allocation ratio N2/N1 = 1

Output: `Noncentrality` parameter δ = `2.8284271`
`Critical t` = `1.9789706`
Df = 126
Sample size group 1 = `64`
Sample size group 2 = `64`
Total `sample size` = `128`
Actual power = 0.8014596

distributions

their cut-offs (type I and II **errors**)

the distance between them (sample and effect **sizes**)

GPower distributions

- distribution based test selection
  - Exact Tests (8)
  - t-tests (11) → `reference`
  - z-tests (2)
  - χ^2-tests (7)
  - F-tests (16)
- focus on the density functions
- design based test selection
  - correlation & regression (15)
  - means (19) → `reference`
  - proportions (8)
  - variances (2)
- focus on the type of parameters

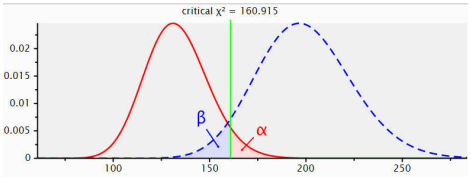

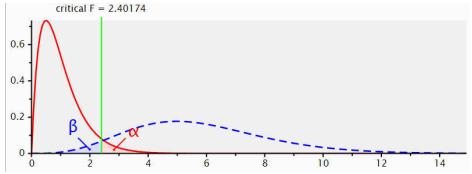

Ho and Ha distributions

- `Ho` acts as *benchmark* → eg., no difference
  - `Ho ~ t(0,df)` *cutoff* using α,
  - reject `Ho` if test returns `implausible` value
- `Ha` acts as *truth* → eg., difference of .5 SD
  - `Ha ~ t(ncp,df)`
  - `ncp` as violation of `Ho` → shift (location/shape)
-

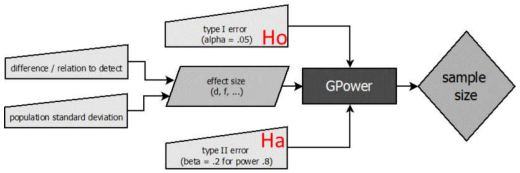

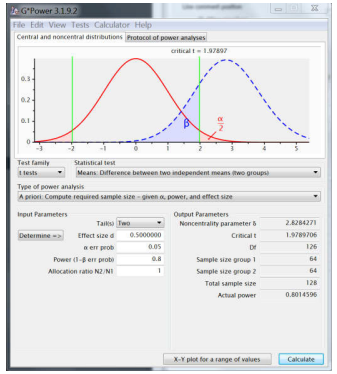

non-centrality: Ho → Ha

- `ncp` : non-centrality parameter
  - shift between `Ho` and `Ha`
    - assumed `effect size` (target or signal)
    - conditional on `sample size` (information)
  - overlap → power or sample size
    using α on `Ho` and β on `Ha`
- `Ha` is NOT interchangeable with `Ho`
  - absence of evidence $\neq$ evidence of absence
  - equivalence testing (`Ha` for 'no effect')

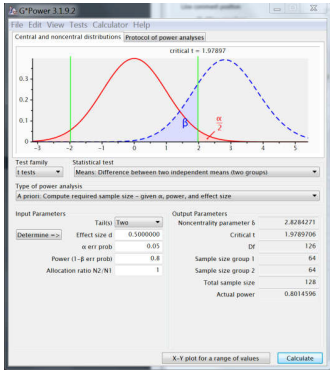

https://icds.shinyapps.io/shinyt/

divide by n perspective on distributions

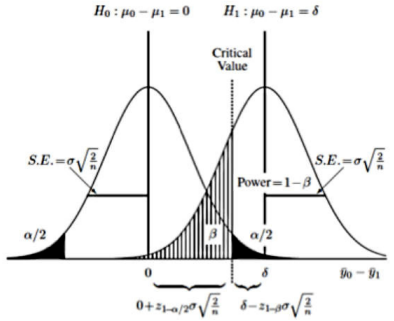

- remember:

$$n = \frac{(Z_{\alpha/2} + Z_\beta)^2 * 2 * \sigma^2}{d^2}$$

$$n = \frac{(-1.96 - 0.84)^2 * 2 * 4^2}{2^2}$$

$$n = 62.79$$

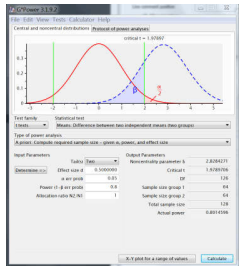

- non-centrality parameter, sample size translates `Ha`
- alternative: sample size changes standard deviation
- https://icds.shinyapps.io/shinyt/

no Ha for estimation

- focus on `estimation`, plausible values of effect, not testing or power
- sample size without type II error β, power, `Ho` or `Ha`
- ~ divide by n perspective but shifted to estimate
- precision analysis → set maximum width confidence interval
  - let E = maximum half width of confidence interval to accept
  - for confidence level $1 - \alpha$
  - $n = z^2_{\alpha/2} * \sigma^2 * 2 / E^2$ (for 2 groups)
- equivalence with statistical testing
  - if 0 (or other reference) outside confidence bounds → significant
- NOT GPower

type I/II error probability

- distribution cut-off's (density → AUC=1)
- decide whether to reject `Ho` **assuming** `Ha`
- two types of error
  - P(infer=`Ha` | truth=`Ho`) = α
  - P(infer=`Ho` | truth=`Ha`) = β
- two types of correct inference
  - P(infer=`Ho` | truth=`Ho`) = $1 - \alpha$
  - P(infer=`Ha` | truth=`Ha`) = $1 - \beta$ → power
- cut-off 'known' `Ho` for statistical test
  - two tailed → both sides informative on `Ho`
  - one tailed → one side not informative on `Ho`

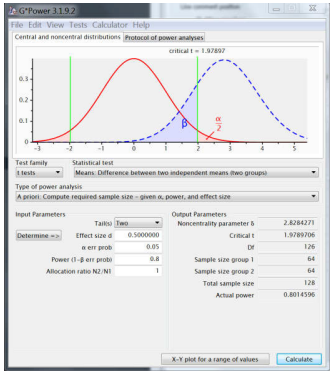

	infer=Ha	infer=Ho	sum
truth=Ho	α	1-α	1
truth=Ha	1-β	β	1

error exercise : create plot

- create plot
  (X-Y plot for range of values)
- plot sample size by type I error
- set plot to 4 curves
  - for power .8 in steps of .05
- set α on x-axis
  - from .01 to .2 in steps of .01
- use effect size .5

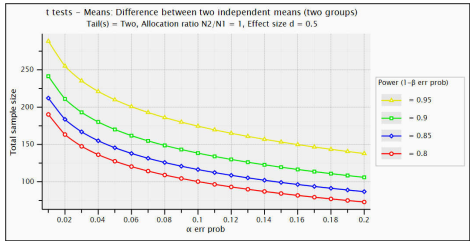

notice Table option

error exercise : interpret plot

- where on the red curve (right)
  type II error = 4 * type I error ?
- when smaller effect size (.25), what changes ?
- switch power and sample size (32 in step of 32)
  what is relation type I and II error ?

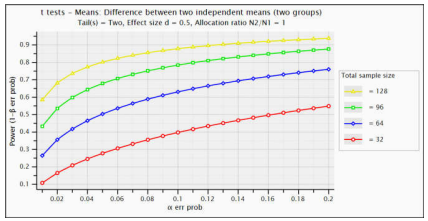

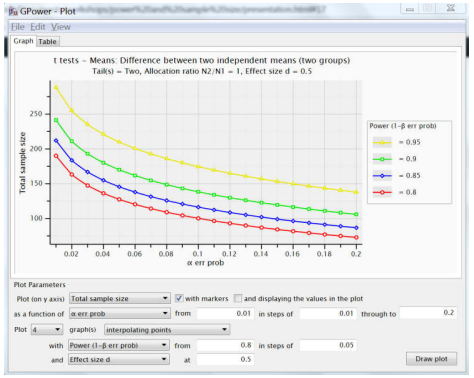

- where on the yellow curve (left)
  type II error = 4 * type I error ?

decide type I/II error probability

- rules of thumb ?
  - α in range .01 - .05 → 1/100 - 1/20
  - β in range .1 to .2 → power = 80% to 90%
- α & β inversely related
  - if $\alpha = 0 \rightarrow$ never reject, no power
  - if power 99% → high α for same sample size
- determine the balance
  - which error you want to avoid most ?
    - cheap aids test ? → avoid type II
    - heavy cancer treatment ? → avoid type I
  - α & β often selected in 1/4 ratio
    type I error is 4 times worse !!

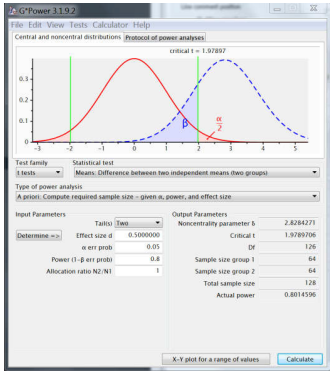

interim analyses: control type I error

- analyze and proceed ? (peeking)
  - multiple testing as data is collected
  - inflates type I error α
- correct α
  - interim analysis specific α_i with overall α under control
- suggested technique: alpha spending
  - plan in advance
  - use O'Brien - Flemming bounds
  - NOT GPower
    - determine with simulation tool: https://icds.shinyapps.io/aspending/
    - commercial packages eg., PASS and other software (eg., ldbounds in R)

for fun: P(effect exists | test says so)

- power → P(test says there is effect | effect exists)
- $P(infer = Ha | truth = Ho) = \alpha$
- $P(infer = Ho | truth = Ha) = \beta$
- $P(infer = Ha | truth = Ha) = power$
- $P(\underline{truth} = Ha | \underline{infer} = Ha) = \frac{P(infer=Ha|truth=Ha) * P(truth=Ha)}{P(infer=Ha)}$ → Bayes Theorem
- $P(truth = Ha | infer = Ha) = \frac{P(infer=Ha|truth=Ha) * P(truth=Ha)}{P(infer=Ha|truth=Ha) * P(truth=Ha) + P(infer=Ha|truth=Ho) * P(truth=Ho)}$
- $P(truth = Ha | infer = Ha) = \frac{power * P(truth=Ha)}{power * P(truth=Ha) + \alpha * P(truth=Ho)}$ → depends on prior probabilities
- IF very low probability model is true, eg., .01 ? → $P(truth = Ha) = .01$
- THEN probability effect exists if test says so is low, in this case only 14% !!
- $P(truth = Ha | infer = Ha) = \frac{.8 * .01}{.8 * .01 + .05 * .99} = .14$

effect sizes

- degree with which a certain phenomenon holds (~ `Ho` is false)
  - part of non-centrality (as is sample size) → shift in GPower
  - signal to noise ratio
    - typically not just the signal, to provide scale
    - eg., difference on scale of pooled standard deviation
  - bigger effect → more easy to detect it (pushing away `Ha`)
- 2 main families of effect sizes → test specific
  - differences `d-family` // association `r-family`
  - transformations, eg., d = .5 → r = .243
    - $d = \frac{2r}{\sqrt{1-r^2}}$; $r = \frac{d}{\sqrt{d^2+4}}$; $d = ln(OR) * \frac{\sqrt{3}}{\pi}$

effect sizes of Cohen

Table 1
ES Indexes and Their Values for Small, Medium, and Large Effects

Test	ES index	Effect size		
		Small	Medium	Large
1. m_A vs. m_B for independent means	$d = \frac{m_A - m_B}{\sigma}$	.20	.50	.80
2. Significance of product–moment r	r	.10	.30	.50
3. r_A vs. r_B for independent rs	$q = z_A - z_B$ where z = Fisher's z	.10	.30	.50
4. P = .5 and the sign test	$g = P - .50$	.05	.15	.25
5. P_A vs. P_B for independent proportions	$h = \phi_A - \phi_B$ where ϕ = arcsine transformation	.20	.50	.80
6. Chi-square for goodness of fit and contingency	$w = \sqrt{\sum_{i=1}^{k} \frac{(P_{1i} - P_{0i})^2}{P_{0i}}}$	.10	.30	.50
7. One-way analysis of variance	$f = \frac{\sigma_m}{\sigma}$	.10	.25	.40
8. Multiple and multiple partial correlation	$f^2 = \frac{R^2}{1 - R^2}$	.02	.15	.35

- famous Cohen conventions
  - beware, just rules of thumb
  - Cohen, J. (1988). Statistical power analysis for the behavioral sciences (2nd ed).
- Cohen, J. (1992). A power primer. Psychological Bulletin, 112, 155–159.

effect sizes of d and r family

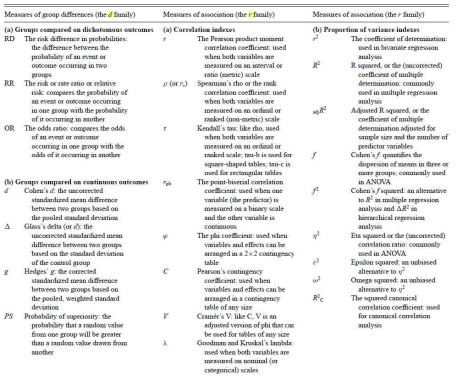

Measures of group differences (the *d* family)

(a) Groups compared on dichotomous outcomes

RD The risk difference in probabilities: the difference between the probability of an event or outcome occurring in two groups

RR The risk or rate ratio or relative risk: compares the probability of an event or outcome occurring in one group with the probability of it occurring in another

OR The odds ratio: compares the odds of an event or outcome occurring in one group with the odds of it occurring in another

(b) Groups compared on continuous outcomes

d Cohen's *d*: the uncorrected standardized mean difference between two groups based on the pooled standard deviation

Δ Glass's delta (or *d*): the uncorrected standardized mean difference between two groups based on the standard deviation of the control group

g Hedges' *g*: the corrected standardized mean difference between two groups based on the pooled, weighted standard deviation

PS Probability of superiority: the probability that a random value from one group will be greater than a random value drawn from another

Measures of association (the *r* family)

(a) Correlation indexes

r The Pearson product moment correlation coefficient: used when both variables are measured on an interval or ratio (metric) scale

r_s (or r_h) Spearman's rho or the rank correlation coefficient: used when both variables are measured on an ordinal or ranked (non-metric) scale

τ Kendall's tau: like rho, used when both variables are measured on an ordinal or ranked scale; tau-b is used for square-shaped tables; tau-c is used for rectangular tables

r_{pb} The point-biserial correlation coefficient: used when one variable (the predictor) is measured on a binary scale and the other variable is continuous

φ The phi coefficient: used when variables and effects can be arranged in a 2×2 contingency table

C Pearson's contingency coefficient: used when variables and effects can be arranged in a contingency table of any size

V Cramér's V: like C, V is an adjusted version of phi that can be used for tables of any size

λ Goodman and Kruskal's lambda: used when both variables are measured on nominal (or categorical) scales

Measures of association (the *r* family)

(b) Proportion of variance indexes

r^2 The coefficient of determination: used in bivariate regression analysis

R^2 R squared, or the (uncorrected) coefficient of multiple determination: commonly used in multiple regression analysis

$adj R^2$ Adjusted R squared, or the coefficient of multiple determination adjusted for sample size and the number of predictor variables

f Cohen's *f*: quantifies the dispersion of means in three or more groups; commonly used in ANOVA

f^2 Cohen's $\sqrt{f}$ squared: an alternative to R^2 in multiple regression analysis and ΔR^2 in hierarchical regression analysis

η^2 Eta squared or the (uncorrected) correlation ratio: commonly used in ANOVA

ε^2 Epsilon squared: an unbiased alternative to η^2

ω^2 Omega squared: an unbiased alternative to η^2

R^2_C The squared canonical correlation coefficient: used for canonical correlation analysis

- Ellis, P. D. (2010). The essential guide to effect sizes: statistical power, meta-analysis, and the interpretation of research results.
- more than 70 different effect sizes... most of them related to each other
- NOT p-value ~ partly effect size, partly sample size or power
  - do not simply compare p-values !
- most important effect size of
  - d: dichotomous - continuous
  - r: correlation - proportion variance

effect sizes in GPower (Determine)

- often very difficult to specify
- GPower offers help with `Determine`
  - difference group means
    0-2 → signal ~ minimally relevant (or expected)
  - standard deviations (sd)
    4 each group → expected noise ~ natural diversity
  - written to Effect Size `d`
    .5 → difference in sd
- Reminder: effect size statistic depends on statistical test

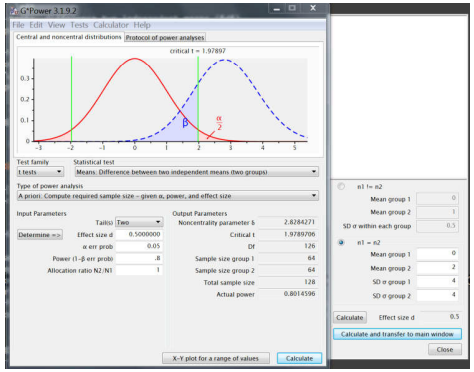

effect size exercise : ingredients cohen d

For the `reference example`:

- change mean values from 0 and 2 to 4 and 6, what changes ?
- change sd values to 2 for each, what changes ?
  - effect size ?
  - total sample size ?
  - non-centrality ?
  - critical t ?
- change sd values to 6 for each, what changes ?

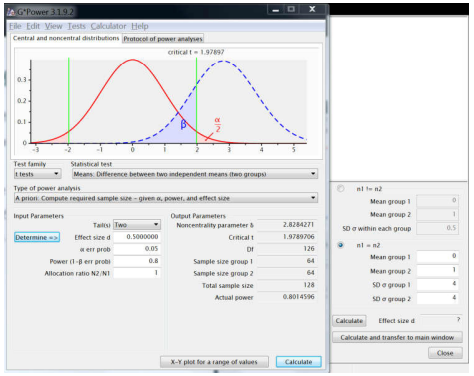

effect size exercise : plot

- plot power by effect size
- set plot to 6 curves
  - for sample sizes, 34 in steps of 34
- set effect sizes on x-axis
  - from .2 to 1.2 in steps of .05
- use α equal to .05
- create plot
  (X-Y plot for range of values)

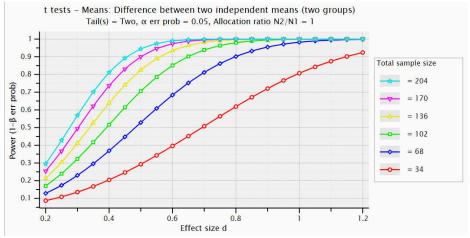

- determine (approximately) the three situations from previous slide on the plot
- how does power change when doubling the effect size, eg., from .5 to 1 ?

effect size exercise : imbalance

For the `reference example`:

- change allocation ratio from 1
  - to 2, .5, 3 and 4, what to conclude ?
    - ratio 2 and .5 ?
    - imbalance + 1 or * 2 ?
- ? no idea why n1 $\neq$ n2

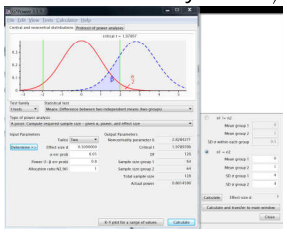

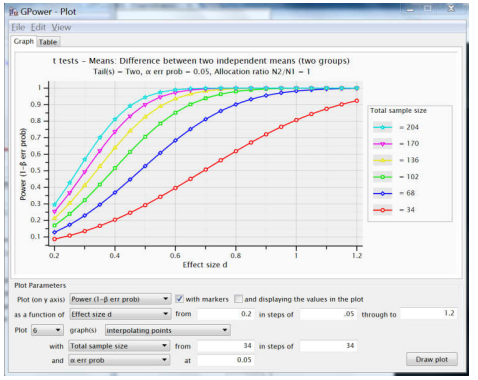

effect sizes, how to determine them in theory

- choice of effect size matters → justify choice
- choice of effect size
  - NOT significant → meaningless, dependent on sample size
  - realistic (eg., previously observed effect) → replicate
  - important (eg., minimally relevant effect)
- use `Determine` to get started (check the manual)
  - for independent t-test → means and standard deviations
  - possible alternative is to use variance explained, eg., 1 versus 16
    - with one-way ANOVA (f=.25 instead of d=.5)
    - with linear regression (f^2=.0625 instead of d=.5)
    - https://www.psychometrica.de/effect_size.html#transform

effect sizes, how to determine them in practice

- experts / patients → use if possible → importance
- literature (earlier study / systematic review) → realistic
- pilot → guestimate dispersion estimate, but very small sample size
- internal pilot → stopping rule (sequential/conditional)
- turn to Cohen → use if everything else fails (rules of thumb)
- guestimate the input parameters, what can you do ?
  - sd from assumed range / 6 assuming normal distribution
  - sd for proportions (& percentages) at conservative .5
  - sd from control, assume treatment the same
  -

relation samples & effect size, errors I & II

- building blocks:
  - sample size (n)
  - effect size (Δ)
  - alpha (α)
  - power ($1 - \beta$)
- each parameter conditional on others
- GPower → type of power analysis
  - Apriori: n ~ α, `power`, Δ
  - Post Hoc: `power` ~ α, n, Δ
  - Compromise: `power`, α ~ β / α, Δ, n
  - Criterion: α ~ `power`, Δ, n
  - Sensitivity: Δ ~ α, `power`, n

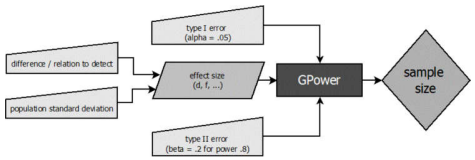

type of power analysis exercise

- for given example, step through...
  - retrieve power given n, α and Δ
  - [1] for power .8, take half the sample size, how does Δ change ?
  - [2] set β/α ratio to 4, what is α & β ? what is the critical value ?
  - [3] keep β/α ratio to 4 for effect size .5, what is α & β ? critical value ?

getting your hands dirty

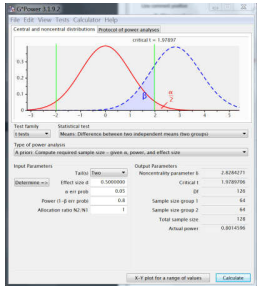

```
# calculator
m1=0;m2=2;s1=4;s2=4
alpha=.025;N=128
var=.5*s1^2+.5*s2^2
d=abs(m1-m2)/sqrt(2*var)
d=d*sqrt(N/2)
tc=tinv(1-alpha,N-1)
power=1-nctcdf(tc,N-1,d)
```

- in R, assuming normality
  - qt → get quantile on Ho ($Z_{1-\alpha/2}$)
  - pt → get probability on Ha (non-central)

```
.n <- 64
.df <- 2*.n-2
.ncp <- 2 / (4 * sqrt(2)) * sqrt(.n)
.power <- 1 -
    pt(
        qt(.975,df=.df),
        df=.df, ncp=.ncp
    ) -
    pt( qt(.025,df=.df), df=.df, ncp=.ncp)
round(.power,4)

## [1] 0.8015
```


PART II: moving beyong independent t-test

GPower examples beyond the independent t-test

- so far, comparing two independent means
- selected topics with small exercises
  - dependent instead of independent
  - non-parametric instead of assuming normality
  - relations instead of groups (regression)
  - correlations
  - proportions, dependent and independent
  - more than 2 groups (compare jointly, pairwise, focused)
  - more than 1 predictor
  - repeated measures
- GPower manual 27 tests: effect size, non-centrality parameter and example !!

dependence between groups

- if 2 dependent groups (eg., before/after treatment) → account for correlations
- matched pairs (t-test / means, difference 2 dependent means)
- use `reference example`
  - [1] use correlation .5 to compare (effect size, ncp, n)
  - [2] how many observations if no correlation exists (`reference example`) ?
  - [3] difference in sample size for correlation .875 ?
  - [4] set original sample size (n=64*2) and effect size (dz=.5), compare ?

non-parametric distribution

- expect non-normally distributed residuals, avoid normality assumption
- only considers ranks or uses permutations → price is efficiency
- avoid when possible, eg., transformations
- two groups → Wilcoxon-Mann-Whitney (t-test / means, diff. 2 indep. means)
- use `reference example`
  - [1] how about n ? compared to parametric → what is % loss efficiency ?
  - [2] change parent distribution to 'min ARE' ? what now ?

relations instead of group differences

- differences between groups → relation observations & categorization
- example → d = .5 → r = .243
- note: slope $\beta = r * \sigma_y / \sigma_x$
- regression coefficient (t-test / regression, one group size of slope)
- sample size for comparing the slope Ha with 0 (=Ho)
  - [1] determine slope (β, with σ_y = 4.12 and σ_x = .5)
  - [2] calculate sample size
  - [3] what if σ_x (predictor values) or σ_y (effect and error) increase ?

relations: a variance perspective

- between and within group variance → relation observations & categorization
- regression coefficient (t-test / regression, fixed model single regression coef)
- use `reference example`, regression style
  - variance within 4^2 and between 1^2, totaling σ_y^2 = 17
  - [1] calculate sample size, compare effect sizes ?
  - [2] what if also other predictors in the model ?
- note: $f^2 = R^2/(1 - R^2)$

more groups to compare, 4 cases

- simple example: assume one control, and two treatments
- if more than two groups, several options
  - test whether at least one differs → omnibus F-test (variances)
  - test whether all differ from eachother → pairwise comparisons
  - test whether selected pairs differ → contrast (t-test)
  - test whether linear combinations of pairs differ → contrasts (t-tests)
    eg., control versus each of the average of treatments

- if multiple tests → inflation of type I error (α)
  - correct α or p-value, eg., using Bonferroni
  - make more tentative inferences

F-test statistic

- multiple groups → not one effect size `d`
- F-test statistic & effect size `f`
- `f` is ratio of variances $\sigma^2_{between}/\sigma^2_{within}$
- example: one control and two treatments
  - `reference example` + 1 group
  - within group observations normally distributed
  - means C=0, T1=2 and T2=4
  - sd for all groups (C,T1,T2) = 4

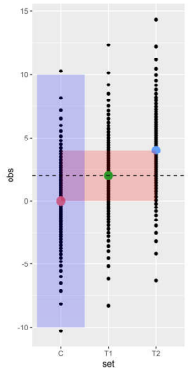

more groups: omnibus

- for one control and two treatments → test that at least one differs
- one-way Anova (F-test / Means, ANOVA - fixed effects, omnibus, one way)
- effect size f, with numerator/denominator df (derived from η^2)
- start from `reference example`,
  - [1] what is the sample size ? ncp ? critical F ? does size matter ?
  - [2] set extra group, either mean 1 or 4, what are the effect / sample sizes ? and with the mean of middle group away from global mean ?
  - [4] derive effect size with `variance` between .5 and within 3, 1 and 6 ?

more groups: pairwise

- assume one control, and two treatments
  - interested in all three pairwise comparisons → maybe Tukey
    - typically run aposteriori, after omnibus shows effect
  - use t-test with correction of α for multiple testing
- apply Bonferroni correction for original 3 group example
  - [1] resulting sample size for three tests ?
  - [2] what if biggest difference is ignored, sample size ?
  - [3] with original 64 sized groups, what is the power ?

sample size calculation benefit from focus

- better to focus during the design on specific questions
  - only consider the main comparisons in focus (eg. primary endpoints)
    - only interested in comparing two treatments → t-test
  - only consider smallest of relevant effects, largest sample size
  - set up contrasts (next slide)
- sample size calculations (design)
  - not necessarily equivalent to statistics
  - requires justification to convince yourself and/or reviewers
- example:
  - statistics: group difference evolution 4 repeated measurements → mixed model
  - power: difference treatment and control last time point → t-test

more groups: contrasts

- assume one control and two treatments
  - set up 2 contrasts for T1 - C and T2 - C
  - set up 1 contrast for average(T1,T2) - C
- each contrast requires 1 degree of freedom
- each contrast combines a specific number of levels
- effect sizes for planned comparisons must be calculated !!
  - contrasts (linear combination)
  - standard deviation of contrasts

$$\sigma_{contrast} = \frac{|\sum \mu_i * c_i|}{\sqrt{N \sum_i^k c_i^2 / n_i}}$$

with group means μ_i, pre-specified coefficients c_i, sample sizes n_i and total sample size N

more groups: contrasts exercise

- one-way ANOVA (F-test / Means, ANOVA-fixed effects,special,main,interaction)
- obtain effect sizes for contrasts (assume equally sized for convenience)
  - $\sigma_{contrast}$ T1-C: $\frac{(-1*0+1*2+0*4)}{\sqrt{(2*((-1)^2+1^2+0^2))}} = 1$; T2-C: $= 2$; (T1+T2)/2-C: $= 1.4142$
  - with σ = 4 → ratio of variances for effect sizes f .25, .5, .3536
- sample size for each contrast, each 1 df and 2 groups
  - [1] contrasts nrs. 1 or 2
  - [2] contrasts nrs. 1 AND 2
  - [3] contrasts nr. 3

multiple factors

- multiway ANOVA (F-test / Means, ANOVA-fixed effects,special,main,interaction)
- multiple main effects and interaction effects
  - interaction: group specific difference between groups
    - degrees of freedom (#A-1)*(#B-1)
  - main effects: if no interaction (#X-1)
  - get effect sizes for two way anova
    https://icds.shinyapps.io/effectsizes/
- sample size for `reference example`, assume a second predictor is trivial
  - [1] what is partial η^2 ?
  - [2] sample size ?

dependence within groups (repeated)

- if repeated measures → account for correlations
- repeated measures (F-test / Means, repeated measures...)
- 3 main types
  - within: like dependent t-test for 2 or more measurements
  - between: use of multiple measurements per group
  - interaction: difference of change over groups
- correlation within subject (unit)
  - informative on group differences within subject
  - redundancy for between group differences

repeated measures within

- possible to have only 1 group (within subject comparison)
- use effect size f = .25 (1/16 explained versus unexplained)
  - [1] use zero correlation to compare with sample size independent t-test
  - [2] for one group use correlation .5, compare sample size dependent t-test
  - [3] double number of groups to 2
  - [4] double number of measurements to 4 (correlation 0 and .5), impact ?

repeated measures between

- use effect size f = .25 (1/16 for variance or 2/4 for means)
  - [1] use correlation 0 and .5 with 2 groups and 2 measurements, sample size ?
  - [2] for correlation .5, compare 2 or 4 measurements, sample size ?
  - [3] double number of groups to 2

repeated measures within x between

- SPSS idiosyncrasies: https://www.youtube.com/watch?v=CEQUNYg80Y0
- use effect size f = .25 (1/16 for variance)
  - [1] use correlation 0, compare 2 groups 2 measurements with rep. between ?
  - [2] use correlation 0.5, compare 2 groups 2 measurements with rep. within ?
  - [3] use correlation .5, compare 2 groups and 4 measurements, sample size ?
  - [4] repeat with 4 groups and 4 measurements, sample size ?

correlations

- when comparing two independent correlations
- z-tests / correlation & regressions: 2 indep. Pearson r's
- makes use of Fisher Z transformations → z = .5 * log($\frac{1+r}{1-r}$) → q = z1-z2
- [1] assume correlation coefficients .7844 and .5 effect size & sample size ?
- [2] assume .9844 and .7, effect size & sample size ?
- [3] assume .1 and .3844 effect size & sample size ?

proportions

- comparing two independent proportions → bounded between 0 and 1
- Fisher Exact Test (exact / proportions, difference 2 independent proportions)
- effect sizes in odds ratio, relative risk, difference proportion
  - [1] for odds ratio 2, p2 = .60, what is p1 ?
  - [2] sample size for equal sized, and type I and II .05 and .2 ?
  - [3] sample size when .95 and .8 (difference of .15) and .05 and .2 ?
- treat as if unbounded, ok within .2 - .8, variance is p*(1-p) → maximally .25 !!
  - [4] use t-test for difference of .15

proportions exercise

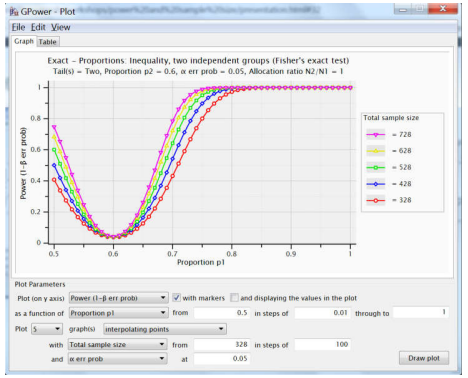

- Fisher Exact Test
  - power over proportions .5 to 1
  - 5 curves, sample sizes 328, 428, 528...
  - type I error .05
  - [1] what happens ?
  - [2] repeat for one-tailed, what is different ?

dependent proportions

- when comparing two dependent proportions
- McNemar test (exact / proportions, difference 2 dependent proportions)
  - include correlations implicitly, discordant pairs → change
  - effect size as odds ratio → ratio of discordance ?!
- assume odds ratio equal to 2, equal sized, type I and II errors .05 and .2, two-way
  - [1] what is the sample size for .25 proportion discordant, and [2] .5, and [3] 1
  - [4] for odds ratio 4 or .25, how the proportion p12 and p21 change ?
  - [5] repeat for third alpha option, and consider total sample size, what happens ?

conclusion: keep it simple, keep it real

- sample size calculation is a design issue, not a statistical one
- building blocks: sample & effect sizes, type I & II errors, each conditional on rest
- effect sizes express the amount of signal compared to the background noise
- complex models imply complex sample size calculations, if at all possible
- GPower deals with not too complex models
  - simplify using a focus, if justifiable → then GPower can get you a long way
  - use more complex specification for more complex sample size calculations
  - leave GPower, simulation is always an option